

PromptFixer · Acceptance test-driven development

Acceptance test results

Run 2026-09-17 19:50:38 UTC · commit f4cea5a (master) · v0.2.0

PASS

113 scenarios (acceptance criteria)	102 passed	0 failed	11 pending, with a stated reason
---	----------------------------------	--------------------------------	--

Every scenario below is a Gherkin acceptance criterion in acceptance/features/, executed by a same-named automated test in acceptance/. Pending scenarios are registered as tests so they appear in every run and cannot be forgotten.

1. Summary by capability

Capability group	Scenarios	Passed	Failed	Pending	Seams
A — Scoring a prompt, with no model involved	12	12	0	0	api
B — Fixing a prompt	10	10	0	0	api, browser
C — The strength contract and the rewrite guard	10	10	0	0	api
D — Learning from the user's own library	6	6	0	0	api
E — The prompt library	7	7	0	0	api
F — The local model	11	3	0	8	api, nightly, desktop, manual
G — Providers and privacy	6	5	0	1	api, nightly
H — The editor and the results interface	14	14	0	0	browser
I — Showing what changed	3	3	0	0	library-seam
J — The CI linter	6	6	0	0	cli
K — The desktop application	6	4	0	2	desktop, nightly, api
L — Configuration, limits and errors	4	4	0	0	api
M — Marking a rewrite and fixing it again	11	11	0	0	api, browser
N — Looking the same every time	7	7	0	0	visual
Total	113	102	0	11	498.3 s of test time

2. Defects found by the scenarios

A8 — "Missing input" was no longer detected for the documented example **FIXED**

DEMO.md and examples.js say the transcript prompt scores a high-severity "references input that is not clearly attached" with Context 24/100. It scored a low "pasted input is not delimited", because the heuristic counted the rest of the instruction sentence after "transcript below" as pasted input.

Fix: server/analyze.js: pasted input starts on a new line or after a colon; the rest of the referencing sentence is instruction.

G2 — An upstream error could carry the API key to the browser **FIXED**

A provider whose error body echoes the key had that body forwarded verbatim, in the error response and inside a successful reply's retry warning. The README promises keys never leave the server.

Fix: server/providers.js redact() blanks every configured key at the source; server/index.js applies it again in the error handler.

D6 — A saved library example returned verbatim was accepted as the rewrite **FIXED**

A parroted example shares enough vocabulary with the prompt to clear the Balanced retention floor, and the leak detector only recognised the `<examples>` wrapper, so another prompt's text was handed back as the user's rewrite.

Fix: `server/metaprompt.js validateRewrite()` receives the examples and treats an exact match with any example as a leak; after two attempts the user gets their original back with a warning naming the library toggle.

K6 / L3 — API errors did not all have the same shape **FIXED**

An unknown `/api` path returned Express's HTML "Cannot GET" page while every other failure is JSON, and the JSON envelope dropped `code/provider` when undefined.

Fix: `server/index.js`: JSON 404 for any unmatched `/api` path through the normal error handler; `code` and `provider` always present (null when unknown).

H12 — Tab count badges were below AA contrast **FIXED**

The count badge on an unselected tab ("Issues 11") used `--text-faint` on `--panel-3`: 4.18:1, below the 4.5:1 the token's own comment promises. Found by measuring every rendered text node in Chromium, not by reading the tokens.

Fix: `src/styles.css`: `.tab .count` uses `--text-dim`.

H13 — The results pane header pushed Save off the page at 1024px wide **FIXED**

Tabs, the history control and Copy/Apply/Save sit in a non-wrapping flex row; at 1024px the Save button ended at 1055px and the whole page scrolled horizontally.

Fix: `src/styles.css`: `.pane-head` wraps.

Observations not acted on

- Ollama shows as configured whenever its default URL is set, even with nothing listening (G1 is stricter than the code; a decision for the three amigos).
- Demo numbers in `DEMO.md` / `examples.js` are stale after the A8 fix (transcript example is now 66/C, not 62/D); the definition of `done` says `DEMO.md` is regenerated when the linter changes.
- `server/data/library.json.tmp-27896` is a crashed-write leftover that is committed to git; it should be removed with `git rm`.
- H8: "download in progress" and "cancel settling" sub-cases need a model file and stay nightly; the other three states are asserted in the browser.

3. Environment

Operating system	Windows_NT 10.0.26200 (x64)
Node.js	v26.8.1
Playwright / Chromium	^1.63.0 / 1.63.0

Electron	^44.1.1
Model provider under test	Scripted OpenAI-compatible stub (no key, no model file); the local model tier is nightly on the GPU machine
Packaged installer	PromptFixer-0.1.0-win-x64.exe · 356 MB · built 2026-09-11 17:47:49 UTC · %LOCALAPPDATA%\PromptFixer\release
Isolation	Each test file boots its own server on an OS-assigned port with private data and model directories

4. Results by acceptance criterion

A — Scoring a prompt, with no model involved

As a person writing prompts I want an instant, deterministic score with concrete findings So that I can improve a prompt before spending any model time on it

ID	Scenario and acceptance criteria	Seam	Result	Time / note
A1	A pasted prompt is scored before anything is contacted • When I paste "write a blog post about our new feature, make it good and professional..." into the editor • Then within a few hundred milliseconds I see a score out of 100 and a letter grade • And I see five category bars: Clarity, Specificity, Context, Format and limits, Structure • And I see a non-empty list of issues • And no request has been made to any provider	api	PASS	323 ms
A2	The same prompt always produces the same score • When the same prompt is analysed twice with the same task type • Then the score, the category breakdown and the ordered issue list are identical	api	PASS	27 ms
A3	Every finding tells the user what to do about it • When a prompt with several defects is analysed • Then every issue has a title, a severity, a category and a suggestion • And issues are ordered from most severe to least	api	PASS	8 ms
A4	A small edit never swings the score wildly • Given a prompt scoring somewhere in the middle • When I add one neutral word to it • Then the overall score moves by at most a few points • And a rewrite that only removes filler words never scores lower than the original	api	PASS	42 ms
A5	A broken short prompt scores below a clean short prompt • Given two prompts of the same length, one clean and one with no verb, format or scope • When both are scored • Then the clean one scores materially higher	api	PASS	17 ms

A6	The task type changes which findings apply • Given the code prompt "fix my login function, it breaks sometimes..." • When it is scored with task type General • Then the "audience not identified" finding is present • When it is scored with task type Code and engineering • Then the "audience not identified" finding is absent • And the "no length or scope limit" finding is low priority rather than a defect	api	PASS	15 ms
A7	Code inside a prompt is not linted as if it were prose • Given a prompt containing a fenced code block with SHOUTED constants and vague-looking identifiers • When it is scored • Then no wording finding cites anything from inside the code • But the structure rules still see the whole prompt	api	PASS	2 ms
A8	A prompt that refers to input it did not attach is a serious finding • Given "Summarize the customer interview transcript below..." with nothing attached • When it is scored • Then there is a high-severity finding that input is referenced but not attached • And its suggestion tells me to add a delimited block • Given the same instruction with the material pasted right below it • When it is scored • Then the missing-input finding is absent and at most a low structure note remains	api	PASS	15 ms
A9	A descriptor-style image prompt is not punished for being one • Given a comma-separated image prompt • When it is scored with task type Image • Then there is no finding for a missing task verb, role or output format	api	PASS	17 ms
A10	A good prompt is told what it got right • Given a well-formed prompt with a role, an output format and a bounded scope • When it is scored • Then it is in the A band • And the panel lists what it does well	api	PASS	2 ms
A11	An oversized prompt is refused politely rather than hanging the app • When a prompt of exactly 60,000 characters is analysed • Then it returns within the interactive budget • When a prompt of 60,001 characters is analysed • Then it is refused with a message naming the length and the limit • And the fix route refuses it the same way	api	PASS	26 ms

A12	An empty editor is an empty state, not an error • When nothing at all is analysed • Then the result is an empty state rather than an error or a zero score presented as a judgement	api	PASS	5 ms
-----	---	-----	------	------

B — Fixing a prompt

As a person with a weak prompt I want the tool to rewrite it and show me exactly what changed and why So that I can trust the rewrite instead of taking it on faith

ID	Scenario and acceptance criteria	Seam	Result	Time / note
B1	The whole journey, end to end • Given the blog prompt that scores 60 with eleven issues • When I fix it at task type Writing and strength Balanced • Then I receive a rewritten prompt, a one-sentence summary, itemised changes with what and why • And any assumptions, questions and techniques the model reports • And a before score and an after score measured with identical settings • And the run details say which provider and model answered and how long it took	api	PASS	1.6 s
B2	The rewrite addresses the problems the user was shown • Given a prompt whose issue list I have read • When I fix it • Then the findings the model was given are exactly the findings I was shown	api	PASS	42 ms
B3	The tool never invents facts about the user's situation • Given a prompt that says "our new feature" without naming it • When the model returns a rewrite containing a marked placeholder • Then the placeholder is scored as a reminder, not a defect, in the after score • But the same placeholder in an original prompt is scored as a defect	api	PASS	44 ms
B4	Failure has a face • Given the provider answers with a server error • When I press Fix • Then I see one sentence naming the provider • And no stack trace, raw JSON or HTTP status appears in the message • And the error says whether retrying is worth it	api	PASS	10 ms
B5	A fix can be cancelled and the machine is handed back • Given the model takes a while to answer • When I cancel the fix mid-generation • Then the server abandons its request to the provider • And no corrective retry is started	api	PASS	1.2 s

B6	A stalled provider gives up in bounded time • Given the provider accepts the connection and never sends headers • When I press Fix • Then the attempt is abandoned at the timeout and marked retryable • Given the provider sends headers and then stalls the body • When I press Fix • Then the attempt is abandoned at the timeout and marked retryable	api	PASS	2.9 s
B7	A truncated reply is reported as a truncated reply • Given the provider stops at its output limit mid-answer • When I press Fix • Then I am told the reply was cut off and that retrying may help • And I am not told the model returned malformed data	api	PASS	10 ms
B8	Nonsense from the client is coerced, never a server error • When options are sent as a string, an array, null, or with unknown extra fields • Then each request succeeds with defaults or fails with a client error • And none of them produces a server error	api	PASS	153 ms
B9	An empty prompt is refused before any model is contacted • When I press Fix with only whitespace in the editor • Then it is refused with an explanation • And no provider call was made	api	PASS	6 ms
B10	The last ten fixes stay comparable • Given I have run twelve fixes in one session • When I use the previous and next controls • Then only the ten most recent are reachable • And the counter reads correctly and clamps at both ends • And the Fixed, Diff, Issues and Changes tabs all follow the selection together	api, browser	PASS	14.7 s

C — The strength contract and the rewrite guard

As a person who chose how much of my text may change I want that choice enforced and explained when the model ignores it So that "Light touch" means what it says

ID	Scenario and acceptance criteria	Seam	Result	Time / note
C1	Strength is a promise about the user's own words • When I fix the same prompt at Light touch, Balanced and Full rebuild • Then the run details state a retention floor of 60%, 30% and none respectively • And the presets offered to the interface list exactly those three strengths	api	PASS	2.0 s

C2	A Light touch rewrite that throws the user's words away is retried, and the faithful attempt is kept • Given the ten-word complaint "I changed the strngth but it still changes a lot" • And the model first returns an entirely different prompt, then a faithful typo fix • When I fix it at Light touch • Then I receive the faithful edit • And the run details show two attempts and no warning • And the retry was told the specific reason the first attempt was rejected	api	PASS	20 ms
C3	When both attempts overshoot, the user is warned and still gets the better one • Given the model returns a different prompt twice • When I fix the complaint at Light touch • Then I still receive a rewrite • And a warning names the strength and the percentage of my words that survived • And the advice is never the strength I am already on	api	PASS	53 ms
C4	Light touch is never asked to do the impossible • Given a prompt whose findings include missing context, audience, format and limits • When I fix it at Light touch • Then the model was shown none of those additive findings • When I fix it at Balanced • Then the model was shown them	api	PASS	49 ms
C5	Keeping every word and bolting scaffolding onto it is also a violation • Given the model keeps every word and appends a role line and three sections, twice • When I fix the complaint at Light touch • Then it is retried once • And the warning says the model added too much, not that it removed too much	api	PASS	23 ms
C6	Full rebuild is genuinely unbounded • Given the model returns an entirely different prompt • When I fix the complaint at Full rebuild • Then I receive it with no warning of any kind • And only one attempt was made	api	PASS	29 ms
C7	The tool's own instructions never reach the user • Given the model pastes its instructions after my text, twice • When I fix the complaint at Light touch • Then the boilerplate is stripped from the result • And every line I wrote survives • And I am warned to check the result carefully • And the output is never empty	api	PASS	72 ms

C8	A user's own finding-shaped text is not mistaken for our boilerplate • Given my prompt legitimately contains bracketed severity tags like a linter report • And the model returns it lightly edited • When I fix it at Light touch • Then I get my text back with no leak warning	api	PASS	24 ms
C9	When two attempts both fall short, the choice between them is predictable • Given the first attempt diverges and the second leaks instructions • When I fix at Light touch • Then the divergent attempt is kept, with a retention warning • Given the first attempt leaks instructions and the second diverges • When I fix at Light touch • Then the divergent attempt is kept, with a retention warning	api	PASS	80 ms
C10	A failed corrective retry never becomes a server error • Given the first attempt overshoots and the retry fails upstream • When I fix at Light touch • Then I receive the first attempt • And the warning also says the retry failed and why	api	PASS	36 ms

D — Learning from the user's own library

As a person who has saved rewrites I liked I want future fixes to match the amount of change I kept before So that the tool adapts to me without any configuration

ID	Scenario and acceptance criteria	Seam	Result	Time / note
D1	The tool learns how much change this user actually likes • Given I have saved a rewrite of a similar blog prompt that scored higher than its original • When I fix a new blog prompt • Then the model was shown that before/after pair as an example • And the run details say one example was used	api	PASS	2.1 s
D2	Only examples worth imitating are used • Given a saved entry whose rewrite scored no better than its original • When I fix a similar prompt • Then that entry was not shown to the model • Given the prompt I am fixing is itself already in the library with a higher after score • When I fix it • Then it was not shown to the model as an example of itself	api	PASS	85 ms

D3	The same task type is preferred • Given two saved entries of equal resemblance, one with task type Writing and one with Code • When I fix a similar prompt at task type Writing and only one example fits • Then the Writing entry is the one shown	api	PASS	60 ms
D4	The user can turn it off, and it never costs anything • Given a saved entry that would normally be used • When I fix a similar prompt with the library toggle off • Then no examples were shown and the run details say zero were used • And the lookup made no network request beyond the model call itself	api	PASS	46 ms
D5	A retry sees the same examples as the first attempt • Given a saved entry that will be used • And the model overshoots on the first attempt • When I fix at Light touch • Then both attempts were shown the same example	api	PASS	47 ms
D6	A rewrite that parrots an example back is caught • Given a saved entry that will be used • And the model returns the example's text as its answer, twice • When I fix a similar prompt • Then it is treated as leaked boilerplate rather than returned as my rewrite	api	PASS	53 ms

E — The prompt library

As a person who fixes prompts regularly I want to keep the ones I liked, find them again, and move them between machines So that my work is never lost and never silently overwritten

ID	Scenario and acceptance criteria	Seam	Result	Time / note
E1	Save it, find it, reopen it, delete it • When I save a completed fix • Then it appears in the library with its scores, provider and model • When I search for a word from the original • Then only matching entries are listed • When I delete it • Then it is gone immediately	api	PASS	465 ms
E2	A hand-edited library file loads anyway • Given the library file contains a null entry, a number where a string belongs, an entry with no id and one valid entry • When I open the library • Then the valid entry is listed and the rest are dropped silently • Given the library file is not valid JSON at all • When I open the library • Then it is empty rather than an error • And the broken file has been moved aside, not deleted	api	PASS	31 ms

E3	A crash during a save cannot corrupt the library • When twenty saves race each other • Then the library file is valid JSON at every moment it is read • And every save is present afterwards • And no temporary file is left behind	api	PASS	102 ms
E4	The library is bounded and drops the oldest first • Given the library holds 500 entries • When I save one more • Then the newest is present, the oldest is gone, and the count is still 500	api	PASS	64 ms
E5	A library moves between machines • When I export the library • Then I get a JSON file named with today's date • When I import that file into a library that already has some of the entries • Then the unknown entries are added and the known ones are skipped, not overwritten • And malformed entries are dropped • And the reply counts what was imported and skipped • When I import the same file again • Then nothing is added	api	PASS	89 ms
E6	Importing the wrong file says so in plain words • When I import a JSON object that is not a library • Then I am told what shape was expected • And the library is unchanged	api	PASS	87 ms
E7	The library lives where the documentation says it lives • Given the data directory override is set • When I save an entry • Then the library file is written inside that directory • And nothing is written to the project's own data folder	api	PASS	27 ms

F — The local model

As a person who wants nothing to leave my machine I want the built-in model to download, load and run without configuration So that the app works offline and I always know its state

ID	Scenario and acceptance criteria	Seam	Result	Time / note
F1	A first-run user with no model is told exactly what to do • Given the local provider is selected and no model is on disk • When I open the app • Then the status says the model is missing and names it with its size • And the catalogue shows every tier and whether it is downloaded • When I press Fix anyway • Then I am told the model is not downloaded, not a generic failure • And linting still works	api	PASS	437 ms

F2	The download is visible, resumable and cancellable • When I start the download • Then I see bytes and a percentage • When I cancel it • Then the state is "missing" with no error and the Download button is back • When I interrupt and restart it • Then it resumes rather than starting from zero	nightly	PENDING	nightly on the GPU machine: needs a model file
F3	A download that produced no usable file says so • Given a download that completes but leaves a file of the wrong size • When it finishes • Then the state is an error with a message, not ready	nightly	PENDING	nightly on the GPU machine: needs a model file
F4	The machine picks a sensible model on its own • Given an explicit tier "lite" in configuration • When I open the app • Then the selected tier is lite regardless of this machine's memory	api	PASS	2.1 s
F5	Switching models is safe at any moment • When I select an unrecognised tier • Then it is rejected with an explanation • When I select a valid tier that is not downloaded • Then the selection applies and the status says it is missing	api	PASS	30 ms
F6	Loading is never confused about which model it is loading • Given a load in flight for one tier • When a request arrives for a different tier • Then it is not satisfied by the load in flight	nightly	PENDING	nightly on the GPU machine: needs a model file
F7	The GPU is used when it exists and its absence is not a failure • When the model loads on this machine • Then the status names the backend, and on CPU the app still works	nightly	PENDING	nightly on the GPU machine: needs a model file
F8	A prompt too long for the model is refused with advice, not truncated • Given a prompt exceeding the model's context • When I press Fix • Then I am told it is too long for the local model and offered alternatives before generation starts	nightly	PENDING	nightly on the GPU machine: needs a model file
F9	The result is always structurally valid • When the smallest model generates on a difficult prompt • Then I get a structured result or a clear error, never a malformed one	nightly	PENDING	nightly on the GPU machine: needs a model file
F10	Quitting releases the model • Given a loaded model • When I close the window • Then the model is unloaded and the process exits within a bounded time	desktop, nightly	PENDING	desktop tier, nightly

F11	Nothing leaves the machine • Given the network is physically disconnected • When I lint, fix, save and export • Then all four succeed	manual	PENDING	manual charter: physically disconnected network
-----	---	--------	---------	---

G — Providers and privacy

*As a person who may have put an API key in configuration I want it used only where I intended and never exposed
So that switching providers is safe*

ID	Scenario and acceptance criteria	Seam	Result	Time / note
G1	The app offers only providers that can actually work • Given only the OpenAI-compatible endpoint is configured • When I read the configuration • Then that provider is marked configured and every other one is not • And the configured default provider is what the interface selects • Given nothing at all is configured • Then the default provider is local	api	PASS	4.3 s
G2	Keys never reach the browser • When every route is exercised, including failures • Then no response body contains the API key	api	PASS	265 ms
G3	A client cannot redirect a keyed provider • Given a request that supplies its own endpoint address for the keyed provider • When I fix a prompt • Then the request still goes to the configured endpoint • Given a request that supplies a non-http address for a key-less endpoint • Then it is refused	api	PASS	19 ms
G4	Model lists are live where possible and honest where not • Given the provider answers with its model list • When I ask for models • Then I get the live list, labelled live • Given the provider rejects the key • When I ask for models • Then I am told the key was rejected, not shown a fallback list • Given a provider with no key configured • When I ask for models • Then I am told the key is missing instead of a guess	api	PASS	44 ms
G5	The journey is the same on every provider • When the same prompt is fixed on the local provider and the compatible one • Then the shape of the result and the guard behaviour are identical	api, nightly	PENDING	nightly: needs the local model

G6	The API is reachable only from this machine • When a request arrives with a foreign origin • Then it is not granted cross-origin access • When a request arrives from a localhost origin • Then it is	api	PASS	22 ms
----	---	-----	------	-------

H — The editor and the results interface

As a first-time user I want the interface to explain itself and never lie about what it is showing So that I can trust the score on screen

ID	Scenario and acceptance criteria	Seam	Result	Time / note
H1	The first thirty seconds • When I open the app for the first time • Then I see an empty editor with an explanatory placeholder, an example gallery and a sample button • When I click an example • Then the prompt, task type and strength are set and the score and issues appear	browser	PASS	15.4 s
H2	All eight built-in examples load and lint • When I choose each example in turn • Then the prompt, task type and strength are set and every other option is left alone	browser	PASS	4.5 s
H3	A broken example list is an empty gallery, not a broken app • Given the examples cannot be fetched • When I open the app • Then the gallery is absent and everything else works	browser	PASS	4.2 s
H4	Editing after a fix does the honest thing • Given a completed fix on screen • When I edit the prompt • Then the Issues tab shows the live analysis of the new text and the fix result is parked	browser	PASS	5.9 s
H5	A deliberately pinned older result keeps its tab • Given I have stepped back to an earlier fix • When I continue to edit • Then the pinned result keeps the Issues tab	browser	PASS	5.3 s
H6	Apply and undo • When I press Apply • Then the fixed text replaces the editor content, the score re-lints and Undo appears • When I press Undo • Then the original text returns and Undo disappears • When I apply and then edit the applied text • Then Undo is no longer offered	browser	PASS	5.8 s
H7	Copy puts the rewrite on the clipboard • When I press Copy • Then the clipboard holds the fixed prompt exactly as displayed	browser	PASS	5.3 s

H8	Fix is disabled when it cannot succeed, and says why • Given each of: empty editor, no provider, model missing, download in progress, cancel settling • Then the Fix button is disabled with a visible reason	browser	PASS	4.8 s
H9	A remembered model choice is reconciled with the server • Given I previously chose a tier and the server is on another • When I open the app • Then my choice is sent to the server, not merely displayed	browser	PASS	3.7 s
H10	Warnings read like sentences a person wrote • Given each kind of guard warning • Then the headline matches the kind and an unknown kind shows the raw text	browser	PASS	16.5 s
H11	The app is usable from the keyboard • When I tab through the interface • Then every control is reachable in a sensible order • And Ctrl+Enter runs a fix, Enter on a nested Delete deletes, Escape closes the drawer	browser	PASS	4.6 s
H12	Text meets contrast requirements • Then every foreground/background pairing meets AA contrast	browser	PASS	4.3 s
H13	The window works at a small size • Given a 1024 by 720 window • Then nothing is clipped and no pane scrolls horizontally	browser	PASS	3.9 s
H14	Every action without a visible result reports itself • When I save, import, export, copy and apply • Then each shows a confirmation stating the outcome	browser	PASS	5.0 s

I — Showing what changed

As a person reading a rewrite I want to see exactly which words changed So that the guard's claims about "how much survived" are verifiable

ID	Scenario and acceptance criteria	Seam	Result	Time / note
I1	The user can see exactly what changed and nothing is hidden • Given an original and a rewrite • When I view the diff • Then removals and additions are distinguishable at word level • Given identical texts • Then the diff shows them as identical	library-seam	PASS	6 ms
I2	One changed word in a long document is one changed word • Given a prompt of over a thousand words with a single word inserted • When I view the diff • Then exactly one insertion is shown and it renders quickly • Given thousands of lines with one line changed • Then the diff finds the one line quickly	library-seam	PASS	50 ms

I3	A region too large to align is reported, not hidden • Given a change too large to align word by word • When I view the diff • Then the region is reported as replaced wholesale, not shown as empty	library-seam	PASS	33 ms
----	---	--------------	------	-------

J — The CI linter

As a team keeping prompts in a repository I want to gate them in CI with the same rules the app uses So that a prompt below the bar cannot be merged

ID	Scenario and acceptance criteria	Seam	Result	Time / note
J1	A team can gate its prompts in CI with no server, no model and no key • When I run the linter over a passing set with a minimum score and a severity threshold • Then it exits 0 • When I run it over a set containing a vague prompt • Then it exits 1 and names the offending file	cli	PASS	4.0 s
J2	The gate uses the same rules as the app • When the same text and task type are scored by the linter and by the app • Then the score, categories and issues are identical	cli	PASS	4.8 s
J3	Directories are searched the way the documentation says • Given a tree with .md, .txt and .prompt files and other file types • When I point the linter at the directory • Then exactly the three extensions are linted, recursively	cli	PASS	2.0 s
J4	A file can declare its own task type • Given a file whose first line is the task-type comment • When it is linted • Then that task type is used and the comment does not affect the score	cli	PASS	6.2 s
J5	Machine-readable output is machine-readable • When I run with --json • Then stdout is exactly one parseable object with every file's score, categories and issues • When I run with --quiet • Then only the summary line is printed	cli	PASS	4.7 s
J6	A misconfigured job is distinguishable from failing prompts • When I run with no paths, an unknown flag, an unknown task type, or a missing path • Then each exits 2 with usage on stderr	cli	PASS	13.1 s

K — The desktop application

As a person who installed PromptFixer I want it to behave like a well-mannered desktop app So that installing, quitting and configuring it hold no surprises

ID	Scenario and acceptance criteria	Seam	Result	Time / note
K1	Install, run, uninstall, reinstall • Given the built installer on a clean machine • When I install, run, uninstall and reinstall • Then the model and the library survive	desktop, nightly	PENDING	desktop tier, nightly: needs the packaged installer on a clean machine
K2	The window opens our interface and nothing else • When a same-origin navigation occurs • Then it proceeds • When a file is dropped, an external link followed, or a lookalike origin navigated to • Then each is blocked in the window	desktop, nightly	PASS	9.3 s
K3	Quitting is deterministic • Given a loaded model • When I quit • Then the quit waits for unload, and a failing or hung unload still quits in bounded time	desktop, nightly	PASS	7.1 s
K4	Configuration is read from where the documentation says • Given the packaged app • Then the per-user .env overrides the desktop defaults and the working-directory .env is ignored • Given the app run from source • Then the project .env wins over the per-user one, and a shell variable wins over both	desktop, nightly	PENDING	desktop tier, nightly: packaged vs source .env matrix (logic covered by electron/env.test.js)
K5	The app finds a free port and its interface finds the app • When two copies of the server start at once • Then both get a port and both answer	api	PASS	5.2 s
K6	An unknown API route is a not-found, not the application shell • When I request an API path that does not exist • Then I get a not-found response, not the interface HTML	api	PASS	3.2 s

L — Configuration, limits and errors

As a person who did not read the configuration section I want the app to work with no configuration and to enforce the limits it advertises So that nothing surprising happens later

ID	Scenario and acceptance criteria	Seam	Result	Time / note
L1	Everything has a working default • Given no configuration at all • When the server starts • Then it starts, selects the local provider and reports the limits it enforces	api	PASS	4.9 s
L2	Port precedence and the development proxy • Given a port of zero • When the server starts • Then it reports the port it was actually given	api	PASS	17 ms

L3	Errors have one consistent shape • When failures are produced on several routes • Then each carries a message, a code, a retryable flag and where relevant a provider	api	PASS	142 ms
L4	The advertised limits are real • When I walk up to the prompt length limit and one step past it • Then the boundary is accepted and the step past it is refused with the limit named • When I walk up to the library cap and one step past it • Then the cap holds	api	PASS	109 ms

M — Marking a rewrite and fixing it again

As a person reading a rewrite that is nearly right I want to point at the parts I loved and the parts I did not So that the next rewrite keeps the first and changes the second without my starting over

ID	Scenario and acceptance criteria	Seam	Result	Time / note
M1	The model is shown what I marked, and the result says how many marks it was given • Given I marked one passage of the rewrite to keep and one to change • When I fix again with my marks • Then the model was shown the rewrite I marked and both passages, each under its own verdict • And it was still asked to fix my original prompt, with the same findings as before • And the new rewrite has the kept passage word for word and no longer has the other as it was • And the run details say one kept passage and one passage to change were applied, and none ignored • And the before score is still the score of my original prompt • Given a rewrite in which two words of the sentence I kept stand once more, on their own, and there I marked them to change • When I fix again with my marks • Then the model is shown both marks, and told that inside the kept sentence the kept sentence wins • And a rewrite that rewords those words where they stood alone is accepted at the first attempt, with no warning	api	PASS	5.8 s
M2	A rewrite that drops a passage I kept is retried, and the faithful attempt is kept • Given the model first paraphrases the passage I kept and leaves the other as it was, then follows my marks • When I fix again with my marks • Then I receive the attempt that follows them, with no warning • And the run details show two attempts • And the retry was told which passage it had dropped and which it had left as it was, each under its own verdict • And it was shown my marks again	api	PASS	95 ms

M3	When the model ignores my marks twice, I still get a result and I am told what was not honoured • Given the model ignores my marks on both attempts • When I fix again with my marks • Then I still receive a rewrite • And a warning says how many kept passages are missing and how many passages to change are still there • And the run details name those passages • And in the interface the warning has its own headline, and the refined note claims no mark that was ignored	api, browser	PASS	10.7 s
M4	Nonsense marks are coerced, never a server error • When marks arrive as a string, a number, a list, null, or holding things that are not text • Then each request succeeds as an ordinary fix and none of them produces a server error • When every mark points at words that are not in the rewrite • Then it is an ordinary fix: the model is shown no marks and the result mentions none • When a passage is marked twice, or padded with spaces, among marks that point at nothing • Then the model is shown it once, trimmed, and nothing else • When a passage is marked both to keep and to change, and those words stand in only one place in the rewrite • Then the change wins • When the same words stand in two places and are marked both ways • Then the model is shown both marks, and rewording one of the two places honours both • When twenty-five passages are marked the same way, or one mark runs to 2,500 characters • Then the model is shown the first twenty, and the first 2,000 characters, and the result counts twenty • When the marked rewrite is longer than a prompt may be • Then it is refused with a message naming the length and the limit	api	PASS	471 ms
M5	A fix without marks is the fix it always was • When I fix the blog prompt with no marks • Then the model is shown no earlier rewrite and no marks • And the result says nothing about marks • Given I have since fixed again with marks • When I fix the blog prompt once more with no marks • Then the earlier marks have no influence on it • Given a model whose rewrite holds a "<user_feedback>" slot of its own for the prompt's user to fill • When I fix the blog prompt with no marks • Then the rewrite comes back with its slot intact, after one attempt and with no warning • And marking that rewrite and fixing again leaves the slot where it was, again with no warning	api	PASS	103 ms

M6	I mark what I loved and what I did not	browser	PASS	14.3 s
	• Given a completed fix on the Fixed tab • Then both marking buttons are disabled, because nothing is selected • When I select words outside the rewrite, or drag from outside into it • Then both buttons stay disabled • When a selection starts in the rewrite and runs on into a toast • Then both buttons are disabled • When a selection of the last line ends where a triple click leaves it, just outside the rewrite • Then that line can be marked • When I double-click a word of the rewrite and press "Keep it — I loved it" • Then that word is highlighted as kept and the selection is gone • When I drag across a sentence and press "Change it — I didn't like it" • Then that sentence is highlighted as to change • And the bar counts one kept and one to change and offers "Fix again with my marks" • Given a rewrite longer than the window, scrolled until the marking bar is stuck above it • When I drag a selection upwards from the rewrite onto the bar, from plain text or from a highlight • Then the selection stays inside the rewrite and both buttons can be pressed • And the bar takes the mouse again when the press ends, even a press that ends without a mouse-up • When twenty passages are kept and I select another • Then "Keep it — I loved it" is disabled and says twenty is the most the model is shown, and the other button is not • When I remove one highlight • Then the passage can be kept			
M7	A mark can be changed and taken back	browser	PASS	6.6 s
	• Given a rewrite with two passages kept and one to change • When I drag across a kept passage again and press "Change it — I didn't like it" • Then the newer verdict replaces the older one • When I click a highlighted passage • Then its highlight is removed and the count follows • When I press Enter on a highlighted passage from the keyboard • Then its highlight is removed too • When I press Ctrl+Enter in the rewrite or on the bar while a mark is showing • Then no fix is started and the mark is still there • When I press "Clear marks" from the keyboard • Then every highlight is gone, and so is the offer to fix again • And the focus is on the rewrite, which shows no ring, and not at the top of the page • When I click in the rewrite and press Ctrl+Enter • Then a fresh fix runs, as it does from anywhere else			

M8	"Fix again with my marks" sends exactly what I marked • Given I kept one passage and marked another to change • When I press "Fix again with my marks" from the keyboard • Then the model was shown exactly those two passages and the rewrite they came from • And I am told it was refined, with the before and after scores • And the new rewrite joins the history as "2 of 2" • And the What changed line says it was refined with my marks, one kept and one changed • And the new rewrite carries no highlights, and the focus is on it, not at the top of the page • And stepping back to the first rewrite does not bring the old highlights back • When I mark the first rewrite again and step forward and back through the history • Then the mark is gone • When I mark it once more and press "Fix prompt" • Then the fresh rewrite carries no highlights, and stepping back to the first one does not bring the mark back • Given the fresh rewrite has two words of the sentence I keep standing once more on their own, and there I mark them to change • When I press "Fix again with my marks" • Then the model was shown one passage to keep and one to change, as the bar counted • And the What changed line says one kept and one changed, with no warning	browser	PASS	5.8 s
M9	The rewrite stays exact while marks are showing • Given highlights of both kinds on the rewrite, one of two lines and one of four lines and more than a hundred characters • Then the rewrite on screen is still the fixed prompt, character for character • And each highlight is a button whose name gives a screen reader the verdict and the whole passage, its line breaks as spaces • When I press Copy • Then the clipboard holds the fixed prompt, with nothing from the marking bar in it	browser	PASS	4.5 s
M10	Marks are readable and fit a small window • Given a 1024 by 720 window with highlights of both kinds and the count showing • Then every piece of text, inside both highlight colours and on the bar, meets AA contrast • And nothing scrolls horizontally and "Fix again with my marks" is on screen	browser	PASS	3.9 s

M11	Fixing again still works after I edited the prompt • Given a completed fix on which I marked a passage • When I type a character in the editor • Then the rewrite is parked behind the edited-prompt notice, and there is nothing to fix again with • When I delete that character again • Then the rewrite is back with my mark on it • When I change the text in the editor and press "Show last fix" • Then the parked rewrite comes back with my mark still on it • When I press "Fix again with my marks" • Then the new rewrite is shown, not parked behind the edited-prompt notice • And the model was asked to refine the prompt that rewrite was made for, not the text now in the editor • And the editor still holds my edit	browser	PASS	4.7 s
-----	--	---------	------	-------

N — Looking the same every time

As a developer changing PromptFixer's interface I want one command to tell me which screens look different, and which element changed So that a change to the look is something I approve, never something a user finds

ID	Scenario and acceptance criteria	Seam	Result	Time / note
N1	Every screen, captured again, matches its approved picture • When every screen is captured again at both window sizes • Then the check passes • And no screen reports a changed, added or removed region at either size	visual	PASS	83.2 s
N2	A rewrite that came back different is reported on the Fixed screen, by name • Given the model now words one line of the rewrite differently • When the Fixed, Issues and Changes screens are captured • Then the check fails • And the Fixed screen reports a change on that one line and names the rewrite as the element that changed • And the Issues and Changes screens, which do not show the rewrite, report nothing	visual	PASS	22.7 s

N3	Marks are part of the picture • When the approved picture of the marked screen is compared with the unmarked one • Then the difference names the marking bar, its "Clear marks" and "Fix again with my marks" buttons and both highlighted passages • And each highlighted passage is a reported region • And nothing in the editor pane, and nothing above the marking bar, is reported • Given both highlights have lost all their paint, and nothing has moved • When the marked screen is captured at both window sizes • Then at each size the check fails and names the two highlighted passages, each with changed pixels behind it, and nothing else • Given the highlight of a kept passage has lost its green • When the marked screen is captured • Then the check fails and names that highlighted passage, and nothing else	visual	PASS	23.3 s
N4	What honestly varies between two runs raises no alarm • Given a model that takes longer to answer and reports other token counts • And a server whose calendar is weeks ahead, so a saved prompt carries another date • When the Fixed, Changes and Library screens are captured • Then the check passes and no screen reports a region	visual	PASS	31.7 s
N5	A screen with no approved picture, or a changed look, fails the check until I approve it • Given approved pictures that lack one screen • And a Fix prompt button that has been restyled • When I run "npm run visual" • Then it fails: the screen without a picture is reported as new, and the other names the Fix prompt button • When I run "npm run visual:approve" • Then it succeeds • When I run "npm run visual" again • Then it passes	visual	PASS	59.7 s
N6	The one command hands back the verdict and leaves nothing running • When I run "npm run visual" on a screen that has not changed • Then it exits 0 • When I run it on a screen that has • Then it exits 1, and a report a CI job can read has been written • When I mistype the name of a screen • Then it exits 2 and lists the screens there are • And after each run the server it started no longer answers and its temporary folders are gone	visual	PASS	18.1 s

N7	Pointed at the wrong place, the check says so and harms nothing • When the check is run with no server to look at • Then it stops with an error that says to use "npm run visual" • Given a PromptFixer that somebody is using, with a prompt saved in its library • When the check is pointed at it • Then it stops with an error, and the saved prompt is still there	visual	PASS	14.2 s
----	---	--------	-------------	--------

Generated by `npm run test:report` from `acceptance\report\out\results.json`. Strategy: `docs/ATDD.md`. Narrative report: `docs/ATDD-REPORT.md`.